

An Implementation of Dictionaries using Randomized Hashing

Avi Weinstock

Jim Olyha

Neil McGlohon

Problem

- Given a large universe U of possible elements, we want to keep track of a set $S \subseteq U$
 - Insert(x) - Add element x to S
 - Delete(x) - Remove element x from S
 - Lookup(x) - Determine if element x is in S
 - Ideally, we want to perform all of these operations in expected constant time.
 - We also don't want to take up more than $O(n)$ space where $n = |S|$

Example: Find name from SSN

- Suppose we want to store a bunch of names and be able to look them up based on the person's Social Security Number
- Naive approach:
 - Make an array of size $m=1,000,000,000$ and store each SSN at the proper index
 - If John Doe's SSN is 123456789, then we set `array[123456789] = "John Doe"`
 - If we are storing much less than 1,000,000,000 people in the array, this is spacially inefficient

A Slightly Better Solution

- Store a linked list of all of the SSNs that we have encountered so far
- Insert each new SSN and Name as a tuple at the beginning of the list - $O(1)$ time
- Search through list when `delete()` or `lookup()` are called.
 - If n elements in list, then these would take $O(n)$ time but only $O(n)$ space (down from $O(m)$)

A Better Solution

- HashTable
 - Efficient Adds - $O(1)$ in expectation
 - Efficient Deletes - $O(1)$ in expectation
 - Efficient Lookups - $O(1)$ in expectation
- HashTable is a Key:Value store
 - Maps a given key to a value

Hash Functions


- Choose where to put the value in the table of elements based on some computation on the key
- Example Hash Function: Use first digit
 - $h(123456789) = 1$
 - $h(213456789) = 2$
 - This obviously would result in an issue...

Hash Functions

- What if we were to add another element:
 - $H(112345678) = 1$
 - This is called a *collision*
 - When there is a collision, we can create a linked list at the hash index

Hash Functions

- Too many collisions?
- Slower lookup() and delete() operations
- Searching over long lists



Writing Good Hash Functions

- Want to minimize the chance of collisions
- We don't want to store much information about location of stored elements
- If we have a table of n elements then the probability of collision between any two elements should be $1/n$.

Random Hashing Schemes: Outline

- Foolish randomized hashing schemes
- Defining a better randomized hashing scheme
 - Universal class of hashing functions
- Bounding probability of collisions to $1/n$
- Given the shown bound, the expected time complexity for each operation (Insert(), Delete(), Lookup()) is $O(1)$

Foolishly Randomized Hashing Scheme

- Pick an index uniformly at random to place an element in the table
- Probability of two elements colliding is $1/n$
 - *Of the n^2 possible choices for the pair of values $(h(u), h(v))$, all are equally likely, and exactly n of these choices results in a collision.⁽¹⁾*
- *What's the problem with this?*

Foolishly Randomized Hashing Scheme

- “Where did I put it?”
- A lookup() or delete() operation would have to search through at most the entire table.
- Well then why not create a hash table that stores the indices of the values that we stored in our original hash table...?
 - Back at square one

Universal Class of Hash Functions

- Set of hash functions such that⁽²⁾:
 - For any pair of elements $u, v \in U$, the probability that a randomly chosen $h \in H$ satisfies that $h(u) = h(v)$ (a collision) is at most $1/n$.
 - Each $h \in H$ can be compactly represented and, for a given $h \in H$ and $u \in U$, we can compute the value $h(u)$ efficiently.

Universal Class of Hash Functions

- If there is a $1/n$ chance of collisions, then we can expect the size of each entry in the table (a linked list) to be 1.
- Therefore the expected time complexity of each operation will be:
 - Insert() - $O(1)$
 - Delete() - $O(1)$
 - Lookup() - $O(1)$

Challenge

- To achieve these expected time complexities we must design good hash functions that fall into the Universal Class of Hash Functions